

# Assignment 3

Group 4

2024-11-08

## Part 1

### Task A

Faults, adjusting clock summer (+02:00) time winter time (+01:00)

handle it by setting UTC as standard to transform to in the ymd\_hms function.

```
power_consum_raw <- read.table("consumption_per_group_aas_hour.csv", sep=";", dec=",", header= T)
colnames(power_consum_raw)
```

```
## [1] "STARTTID"          "SLUTTID"           "KOMMUNE"           "FORBRUKSGRUPPE"
## [5] "VOLUM_KWH"         "ANTALL_MÅLEPUNKT"
```

```
power_consum_df <- power_consum_raw %>%
  dplyr::select(c("STARTTID", "FORBRUKSGRUPPE", "VOLUM_KWH")) %>%
  mutate(STARTTID = ymd_hms(STARTTID, tz = "UTC"))
```

```
summary(power_consum_df)
```

```
##      STARTTID          FORBRUKSGRUPPE      VOLUM_KWH
## Min.   :2021-03-31 22:00:00.00 Length:90231 Min.   : 862.2
## 1st Qu.:2022-03-06 00:00:00.00 Class :character 1st Qu.: 3682.7
## Median :2023-01-13 07:00:00.00 Mode  :character Median : 8751.0
## Mean   :2023-01-13 04:55:16.93      Mean   : 9307.1
## 3rd Qu.:2023-11-22 14:00:00.00      3rd Qu.:12655.9
## Max.   :2024-09-30 21:00:00.00      Max.   :36386.2
```

```
head(power_consum_df)
```

```
##      STARTTID FORBRUKSGRUPPE VOLUM_KWH
## 1 2021-03-31 22:00:00      Forretning 9883.478
## 2 2021-03-31 22:00:00      Industri 2848.522
## 3 2021-03-31 22:00:00      Privat 15463.043
## 4 2021-03-31 23:00:00      Forretning 9867.467
## 5 2021-03-31 23:00:00      Industri 2687.639
## 6 2021-03-31 23:00:00      Privat 14847.819
```

```
power_consum_df_long <- power_consum_df %>%
  pivot_wider(names_from = FORBRUKSGRUPPE, values_from = VOLUM_KWH)
```

```
data_range <- seq(min(power_consum_df_long$STARTTID), max(power_consum_df_long$STARTTID), by = "1 hour")
```

```
last_date_gap <- tail(data_range[!data_range %in% power_consum_df_long$STARTTID], 1)
```

```
power_consum_df_long_cut <- power_consum_df_long %>%
  dplyr::filter(STARTTID >= last_date_gap)
```

```
head(power_consum_df_long_cut)
```

```
## # A tibble: 6 x 4
##   STARTTID          Forretning Industri Privat
##   <dtm>              <dbl>      <dbl> <dbl>
## 1 2021-04-30 22:00:00      8904.    2444. 13809.
## 2 2021-04-30 23:00:00      8934.    2412. 12894.
## 3 2021-05-01 00:00:00      8766.    2464. 12543.
## 4 2021-05-01 01:00:00      8967.    2548. 12502.
## 5 2021-05-01 02:00:00      9067.    2309. 12622.
## 6 2021-05-01 03:00:00      8930.    2334. 12794.
```

```
head(power_consum_df_long_cut)
```

```
## # A tibble: 6 x 4
##   STARTTID          Forretning Industri Privat
##   <dtm>              <dbl>      <dbl> <dbl>
## 1 2021-04-30 22:00:00      8904.    2444. 13809.
## 2 2021-04-30 23:00:00      8934.    2412. 12894.
## 3 2021-05-01 00:00:00      8766.    2464. 12543.
## 4 2021-05-01 01:00:00      8967.    2548. 12502.
## 5 2021-05-01 02:00:00      9067.    2309. 12622.
## 6 2021-05-01 03:00:00      8930.    2334. 12794.
```

```
tail(power_consum_df_long_cut)
```

```
## # A tibble: 6 x 4
##   STARTTID          Forretning Industri Privat
##   <dtm>              <dbl>      <dbl> <dbl>
## 1 2024-09-30 16:00:00      9972.    3281. 15494.
## 2 2024-09-30 17:00:00      9744.    3251. 15659.
## 3 2024-09-30 18:00:00      9294.    2844. 15680.
## 4 2024-09-30 19:00:00      8797.    2624. 15352.
## 5 2024-09-30 20:00:00      8495.    2593. 14576.
## 6 2024-09-30 21:00:00      7894.    2528. 13560.
```

```
# Manually remove the head and tail that don't sum to 1 whole day
```

```
power_consum_df_long_cut_cutDay <- power_consum_df_long_cut %>%
  dplyr::filter(STARTTID > ymd_hms("2021-04-30 23:00:00 UTC")) %>%
  dplyr::filter(STARTTID < ymd_hms("2024-09-30 00:00:00 UTC"))
head(power_consum_df_long_cut_cutDay)
```

```
## # A tibble: 6 x 4
##   STARTTID          Forretning Industri Privat
##   <dtm>              <dbl>      <dbl> <dbl>
## 1 2021-05-01 00:00:00      8766.    2464. 12543.
## 2 2021-05-01 01:00:00      8967.    2548. 12502.
## 3 2021-05-01 02:00:00      9067.    2309. 12622.
## 4 2021-05-01 03:00:00      8930.    2334. 12794.
## 5 2021-05-01 04:00:00      8926.    2174. 13112.
## 6 2021-05-01 05:00:00      8841.    1941. 13587.
```

```
tail(power_consum_df_long_cut_cutDay, 24)
```

```
## # A tibble: 24 x 4
##   STARTTID      Forretning Industri Privat
##   <dtm>          <dbl>    <dbl>  <dbl>
## 1 2024-09-29 00:00:00    7809.    2039. 12014.
## 2 2024-09-29 01:00:00    8003.    2037. 11811.
## 3 2024-09-29 02:00:00    8082.    2034. 11777.
## 4 2024-09-29 03:00:00    8103.    2057. 11791.
## 5 2024-09-29 04:00:00    8435.    2079. 12148.
## 6 2024-09-29 05:00:00    8496.    2001. 12676.
## 7 2024-09-29 06:00:00    8600.    1507. 13679.
## 8 2024-09-29 07:00:00    8095.    1257. 14269.
## 9 2024-09-29 08:00:00    8140.    1255. 14333.
## 10 2024-09-29 09:00:00    9174.    1212. 14255.
## # i 14 more rows
```

```
power_consum_sum_df <- power_consum_df_long_cut_cutDay %>%
  mutate(STARTTID = as.Date(STARTTID)) %>%
  group_by(STARTTID) %>%
  summarise_each(funs(sum))
```

```
## Warning: `summarise_each()` was deprecated in dplyr 0.7.0.
## i Please use `across()` instead.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## Warning: `funs()` was deprecated in dplyr 0.8.0.
## i Please use a list of either functions or lambdas:
##
## # Simple named list: list(mean = mean, median = median)
##
## # Auto named with `tibble::lst()`: tibble::lst(mean, median)
##
## # Using lambdas list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
head(power_consum_sum_df)
```

```
## # A tibble: 6 x 4
##   STARTTID      Forretning Industri Privat
##   <date>          <dbl>    <dbl>  <dbl>
## 1 2021-05-01    205741.    50381. 339827.
## 2 2021-05-02    204183.    50310. 329856.
## 3 2021-05-03    259292.    82488. 329644.
## 4 2021-05-04    270244.    84911. 363652.
## 5 2021-05-05    293115.    91995. 408617.
## 6 2021-05-06    275699.    82163. 380373.
```

Final clean

```
power_consum_df_c <- power_consum_sum_df
```

## Task B

Data is daily.

Data registered on different times, might skew results.

In the excell file the dates are stored in this format '01/01/2020 22:10:00'

Relevant years: 2017-2024:

```
•  
file_path <- getwd()  
excel_files <- list.files(path = file_path, pattern = "*.xlsx", full.names = TRUE)
```

Data formats that might be problematic, some non unique dato. Leap years. Files DATE formatted differently, might cause issues if the read\_excel function does not account for it.

```
metro_aas_df <- excel_files %>%  
  lapply(read_excel) %>%  
  bind_rows() %>%  
  dplyr::select(c("DATO", "LT", "GLOB")) %>%  
  mutate(DATO = as.Date(DATO)) %>%  
  group_by(DATO) %>%  
  distinct(DATO, .keep_all = TRUE) %>%  
  ungroup() %>%  
  complete(DATO = seq(min(DATO), max(DATO), by = 'day')) # Add missing days
```

```
## New names:  
## * `` -> `...30`
```

```
summary(metro_aas_df)
```

##	DATO	LT	GLOB
##	Min. :2017-01-01	Min. :-20.111	Min. : 0.09481
##	1st Qu.:2018-12-12	1st Qu.: 1.086	1st Qu.: 2.14463
##	Median :2020-11-22	Median : 7.049	Median : 7.76017
##	Mean :2020-11-22	Mean : 7.121	Mean :10.33553
##	3rd Qu.:2022-11-02	3rd Qu.:14.142	3rd Qu.:17.39571
##	Max. :2024-10-13	Max. : 24.593	Max. :31.88862
##		NA's :3	NA's :11

```
head(metro_aas_df)
```

```
## # A tibble: 6 x 3  
##   DATO          LT  GLOB  
##   <date>      <dbl> <dbl>  
## 1 2017-01-01    0.543 1.34  
## 2 2017-01-02   -4.44 1.48  
## 3 2017-01-03   -2.37 0.794  
## 4 2017-01-04   -3.18 0.909  
## 5 2017-01-05  -10.8  1.85  
## 6 2017-01-06   -4.18 0.530
```

```
max(metro_aas_df$DATO)
```

```
## [1] "2024-10-13"
```


## Missing data analysis

```
metro_aas_df_c <- metro_aas_df
```

```
ggplot_na_distribution(metro_aas_df_c$LT)
```

### Distribution of Missing Values


Time Series with highlighted missing regions



```
ggplot_na_distribution(metro_aas_df_c$GLOB)
```

## Distribution of Missing Values

Time Series with highlighted missing regions




As seen in the previous graphs, there are very few missing values and they seem evenly spread out, therefore imputations seems to be a good choice to fill in the missing days.

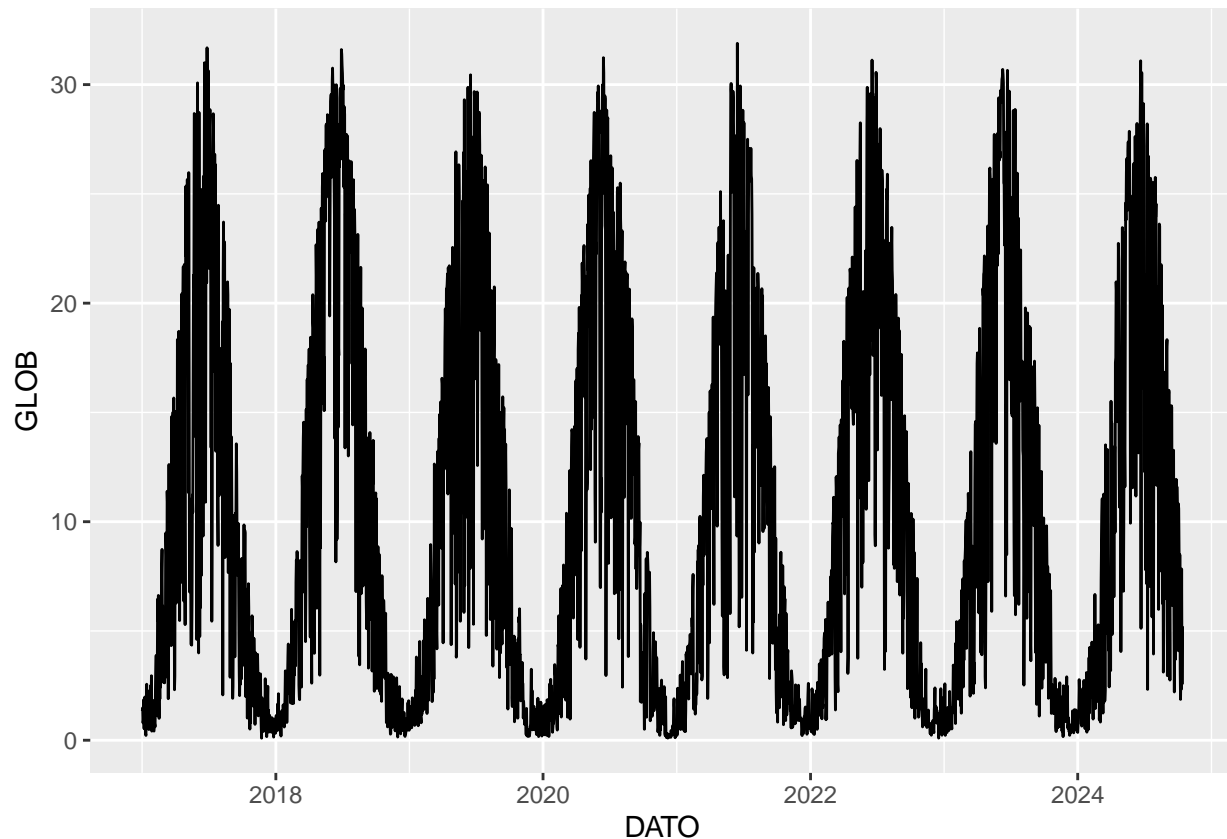
```
metro_aas_df_imputed <- imputeTS::na_ma(metro_aas_df_c , k = 4, weighting = "simple")
head(metro_aas_df_imputed)
```

```
## # A tibble: 6 x 3
##   DATO      LT  GLOB
##   <date>    <dbl> <dbl>
## 1 2017-01-01  0.543  1.34
## 2 2017-01-02 -4.44   1.48
## 3 2017-01-03 -2.37   0.794
## 4 2017-01-04 -3.18   0.909
## 5 2017-01-05 -10.8    1.85
## 6 2017-01-06 -4.18   0.530
```

```
metro_aas_df_imputed %>%
  ggplot(aes(DATO, LT)) +
  geom_line()
```



```
metro_aas_df_imputed %>%  
  ggplot(aes(DATO, GLOB)) +  
  geom_line()
```



### Final cleaned data

```
metro_aas_df_c <- metro_aas_df_imputed
```

## Task C

Find the range of dates for each of the two data sets.

```
power_consum_dateRange <- c(min(power_consum_df_c$STARTTID), max(power_consum_df_c$STARTTID))
metro_ass_dateRange <- c(min(metro_aas_df_c$DATO), max(metro_aas_df_c$DATO))
power_consum_dateRange
```

```
## [1] "2021-05-01" "2024-09-29"
```

```
metro_ass_dateRange
```

```
## [1] "2017-01-01" "2024-10-13"
```

For the longest contiguous range of dates present in both data sets, merge the two data sets based on date.

```
merged_df <- dplyr::inner_join(power_consum_df_c, metro_aas_df_c, by = join_by(STARTTID == DATO)) %>%
  rename(DATO = STARTTID)
```

```
head(merged_df)
```

```
## # A tibble: 6 x 6
```

```
##   DATO      Forretning Industri Privat   LT  GLOB
```



```
##   <date>           <dbl>      <dbl>   <dbl> <dbl> <dbl>
## 1 2021-05-01      205741.    50381.  339827.  6.28 20.9
## 2 2021-05-02      204183.    50310.  329856.  7.45 22.1
## 3 2021-05-03      259292.    82488.  329644.  5.97 16.3
## 4 2021-05-04      270244.    84911.  363652.  5.06  9.37
## 5 2021-05-05      293115.    91995.  408617.  3.36  5.69
## 6 2021-05-06      275699.    82163.  380373.  5.11 17.3
```

Remove data for leap days.

```
merged_df_noLeap <- merged_df %>%
  dplyr::filter(!(month(DATO) == 2 & day(DATO) == 29))

head(merged_df_noLeap)
```

```
## # A tibble: 6 x 6
##   DATO      Forretning Industri  Privat    LT  GLOB
##   <date>      <dbl>      <dbl>   <dbl> <dbl> <dbl>
## 1 2021-05-01    205741.    50381.  339827.  6.28 20.9
## 2 2021-05-02    204183.    50310.  329856.  7.45 22.1
## 3 2021-05-03    259292.    82488.  329644.  5.97 16.3
## 4 2021-05-04    270244.    84911.  363652.  5.06  9.37
## 5 2021-05-05    293115.    91995.  408617.  3.36  5.69
## 6 2021-05-06    275699.    82163.  380373.  5.11 17.3
```

```
merged_data <- merged_df_noLeap
```


## Part 2

### Task D

1. **Data Visualization** Forretning, Industri, Privat are all VOLUME\_KWH

```
# Plot raw data visualization
```

```
merged_data %>%
  pivot_longer(cols = c(Forretning, Industri, Privat, LT, GLOB), names_to = "Measurement", values_to = "Value") +
  ggplot(aes(x = DATO, y = Value)) +
  geom_line() +
  facet_wrap(~Measurement, scales = "free_y") +
  theme_bw() +
  labs(title = "Raw Data Visualization", x = "Date", y = "Value") +
  theme(legend.position = "bottom")
```



```
test_stationarity <- function(data, group_name) {
  print(paste("Stationarity Test for", group_name))
  # Handle NAs (Important: Use na.remove *inside* the function)
  data <- na.remove(data)
  adf.test(data) %>% print() # Augmented Dickey-Fuller Test
  kpss.test(data) %>% print() # KPSS Test
  cat("\n") # Newline for better output formatting
}
```

```
ts_Forr <- ts(merged_data$Forretning, frequency = 365)
test_stationarity(ts_Forr, "Forretning KWH")
```

## 2. Stationarity Tests

```
## [1] "Stationarity Test for Forretning KWH"
##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -2.7124, Lag order = 10, p-value = 0.2768
## alternative hypothesis: stationary
##
## Warning in kpss.test(data): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 0.99076, Truncation lag parameter = 7, p-value = 0.01
```

```

ts_Indu <- ts(merged_data$Industri, frequency = 365)
test_stationarity(ts_Forr, "Industri KWH")

## [1] "Stationarity Test for Industri KWH"
##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -2.7124, Lag order = 10, p-value = 0.2768
## alternative hypothesis: stationary
## Warning in kpss.test(data): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 0.99076, Truncation lag parameter = 7, p-value = 0.01
ts_Priv <- ts(merged_data$Forretning, frequency = 365)
test_stationarity(ts_Priv, "Privat KWH")

## [1] "Stationarity Test for Privat KWH"
##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -2.7124, Lag order = 10, p-value = 0.2768
## alternative hypothesis: stationary
## Warning in kpss.test(data): p-value smaller than printed p-value
##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 0.99076, Truncation lag parameter = 7, p-value = 0.01
ts_temp <- ts(merged_data$LT, frequency = 365)
test_stationarity(ts_temp, "Temperature (LT)")

## [1] "Stationarity Test for Temperature (LT)"
##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -2.4565, Lag order = 10, p-value = 0.3851
## alternative hypothesis: stationary
##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 0.54711, Truncation lag parameter = 7, p-value = 0.03106
ts_glob <- ts(merged_data$GLOB, frequency = 365)
test_stationarity(ts_glob, "Global Irradiation (GLOB)")

```

```
## [1] "Stationarity Test for Global Irradiation (GLOB)"
##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -2.5149, Lag order = 10, p-value = 0.3603
## alternative hypothesis: stationary
##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 0.41001, Truncation lag parameter = 7, p-value = 0.07284
```

```
cor_matrix <- cor(merged_data %>% select(where(is.numeric)), use = "pairwise.complete.obs")
print(cor_matrix)
```


### 3. Cross-Correlation Analysis

```
##          Forretning  Industri  Privat      LT      GLOB
## Forretning 1.0000000  0.9628401  0.8223101 -0.8274985 -0.6222319
## Industri   0.9628401  1.0000000  0.8164847 -0.8050398 -0.6797756
## Privat     0.8223101  0.8164847  1.0000000 -0.9631797 -0.6887294
## LT         -0.8274985 -0.8050398 -0.9631797  1.0000000  0.6716545
## GLOB       -0.6222319 -0.6797756 -0.6887294  0.6716545  1.0000000
```

```
melt(cor_matrix) %>%
  ggplot(aes(Var2, Var1)) +
  geom_tile(aes(fill = value)) +
  geom_text(aes(fill = value, label = round(value, 3))) +
  scale_fill_gradient2() +
  labs(title = "Correlation Matrix", x = "", y = "") +
  theme_minimal()
```

```
## Warning in geom_text(aes(fill = value, label = round(value, 3))): Ignoring
## unknown aesthetics: fill
```


Correlation Matrix

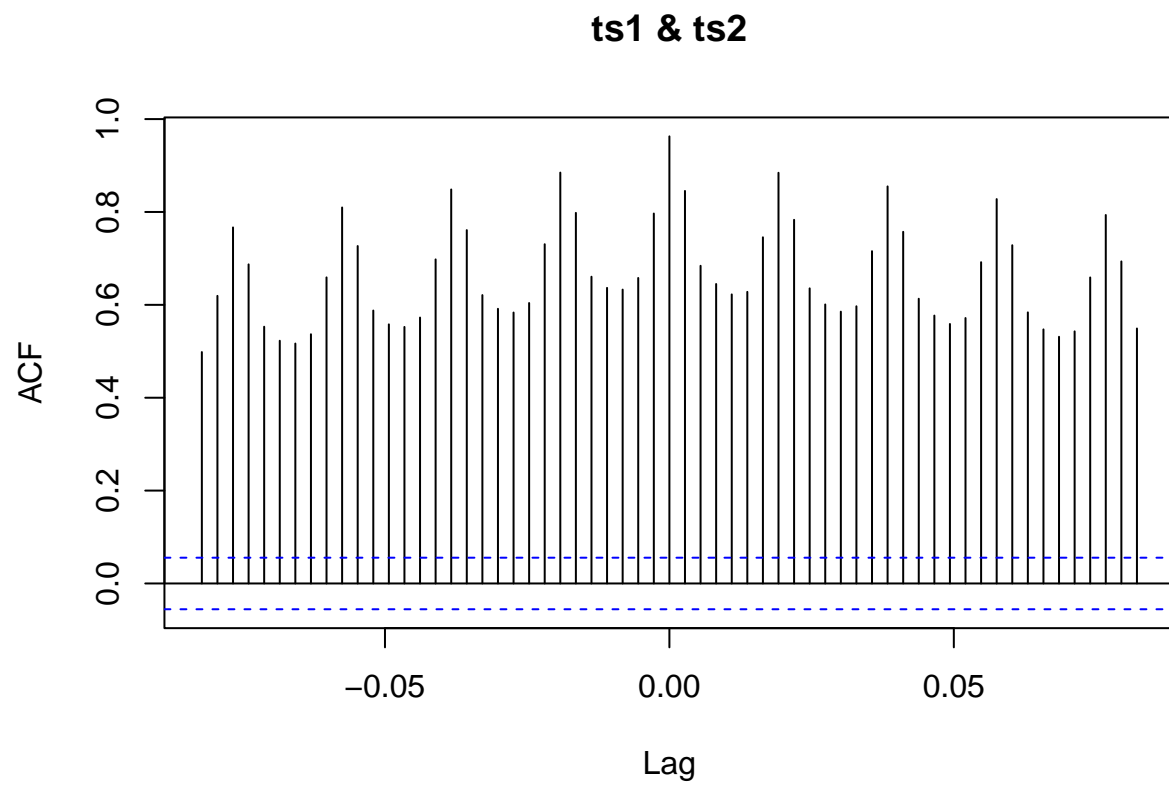


```
numeric_cols <- names(merged_data)[sapply(merged_data, is.numeric) & names(merged_data) != "DATO"]


for (i in 1:(length(numeric_cols) - 1)) {
  for (j in (i + 1):length(numeric_cols)) {
    ts1 <- ts(merged_data[[numeric_cols[i]]], frequency = 365)
    ts2 <- ts(merged_data[[numeric_cols[j]]], frequency = 365)
    ccf_result <- ccf(ts1, ts2, lag.max = 30, na.action = na.contiguous,
                      main = paste("Cross-Correlation:", numeric_cols[i], "vs.", numeric_cols[j]))
    plot(ccf_result)
  }
}
```

### Cross-Correlation: Forretning vs. Industri






### Cross-Correlation: Forretning vs. Privat







**ts1 & ts2**




### Cross-Correlation: Forretning vs. LT




**ts1 & ts2**




### Cross-Correlation: Forretning vs. GLOB




**ts1 & ts2**




### Cross-Correlation: Industri vs. Privat



**ts1 & ts2**




### Cross-Correlation: Industri vs. LT







**ts1 & ts2**




### Cross-Correlation: Industri vs. GLOB




### ts1 & ts2




### Cross-Correlation: Privat vs. LT




**ts1 & ts2**




### Cross-Correlation: Privat vs. GLOB



**ts1 & ts2**




### Cross-Correlation: LT vs. GLOB






## ts1 & ts2




```
acf_pacf_plot <- function(data) {  
  ts_data <- ts(na.remove(data), frequency = 365)  
  
  acf(ts_data, lag.max = 14, main = paste("ACF:", deparse(substitute(data)) , "(Short Term)"))  
  pacf(ts_data, lag.max = 14, main = paste("PACF:", deparse(substitute(data)) , "(Short Term)"))  
  
  acf(ts_data, lag.max = 365*2, main = paste("ACF:", deparse(substitute(data)) , "(Long Term)"))  
  pacf(ts_data, lag.max = 365*2, main = paste("PACF:", deparse(substitute(data)) , "(Long Term)"))  
}  
  
acf_pacf_plot(merged_data$Forretning)
```

ACF: merged\_data\$Forretning




#### 4. Autocorrelation and Partial Autocorrelation Functions (ACF and PACF)


PACF: merged\_data\$Forretning (Short Term)



**ACF: merged\_data\$Forretning (Long Term)**




### PACF: merged\_data\$Forretning (Long Term)




```
acf_pacf_plot(merged_data$Industri)
```


**ACF: merged\_data\$Industri (Short Term)**




**PACF: merged\_data\$Industri (Short Term)**



**ACF: merged\_data\$Industri (Long Term)**




### PACF: merged\_data\$Industri (Long Term)




```
acf_pacf_plot(merged_data$Privat)
```




**ACF: merged\_data\$Privat (Short Term)**




**PACF: merged\_data\$Privat (Short Term)**



**ACF: merged\_data\$Privat (Long Term)**




### PACF: merged\_data\$Privat (Long Term)




```
acf_pacf_plot(merged_data$LT)
```


**ACF: merged\_data\$LT (Short Term)**




**PACF: merged\_data\$LT (Short Term)**



ACF: merged\_data\$LT (Long Term)




### PACF: merged\_data\$LT (Long Term)




```
acf_pacf_plot(merged_data$GLOB)
```




**ACF: merged\_data\$GLOB (Short Term)**




## PACF: merged\_data\$GLOB (Short Term)



ACF: merged\_data\$GLOB (Long Term)



### PACF: merged\_data\$GLOB (Long Term)




### Task E: Seasonal Differencing and ACF/PACF

```
ts_data <- ts(na.remove(merged_data$Privat), frequency = 365)
diff_consumption <- diff(ts_data, lag = 365)

acf(diff_consumption, lag.max = 14, main = paste("ACF:(Seasonally Differenced, Short Term)"))
```


ACF:(Se



Seasonal Differencing of Consumption, Temperature, and Global Irradiation

```
pacf(diff_consumption, lag.max = 14, main = paste("PACF:(Seasonally Differenced, Short Term)"))
```

## PACF:(Seasonally Differenced, Short Term)



*# Repeat for Temperature and Global Irradiation...*

## Task F: STL Decomposition Analysis


```
# Choose STL parameters
stl_parameters <- list(s.window = "periodic", # Seasonal window for yearly data
t.window = 365, # Trend window (yearly smoothing)
l.window = 30, # Low-frequency window (adjust as needed)
robust = TRUE) # Use robust STL (handles outliers better)
ma_order <- 7 # Set the order for the moving average smoother (e.g., 7 for weekly smoothing). You can

# --- Private Electricity Consumption ---
consumption_data <- merged_data$Privat
ts_consumption <- ts(na.remove(consumption_data), frequency = 365)
# Perform STL decomposition
stl_consumption <- stlplus(ts_consumption, period = 365,
                           s.window = stl_parameters$s.window,
                           t.window = stl_parameters$t.window,
                           l.window = stl_parameters$l.window,
                           robust = stl_parameters$robust)
# Smooth the trend-cycle (AFTER STL) using a moving average
smoothed_trendcycle <- ma(stl_consumption$data[, "trend"], order = ma_order) # Using moving average.
# Extract components
seasonal_consumption <- stl_consumption$data[, "seasonal"] # Seasonal component (periodic fluctuations)
remainder_consumption <- stl_consumption$data[, "remainder"] # Remainder component (noise or error)
trendcycle_consumption <- stl_consumption$data[, "trend"]
```

```

if ("cycle" %in% colnames(stl_consumption$data)) {
  trendcycle_consumption <- trendcycle_consumption + stl_consumption$data[, "cycle"]
}
deseasoned_consumption <- ts_consumption - seasonal_consumption # Deseasonalized data
# Plot components
plot(ts_consumption, main = "Private Electricity Consumption") # Original


```



```

plot(seasonal_consumption, main = "Seasonal Component") # Seasonal


```




```
plot(deseasoned_consumption, main = "Deseasoned Consumption") # Deseasoned
```




## Deseasoned Consumption



```
plot(smoothed_trendcycle, main = "Smoothed Trend-Cycle")    # Smoothed trend-cycle
```




```
plot(remainder_consumption, main = "Remainder")           # Remainder
```




```
# --- Global Irradiation ---
ts_glob <- ts(na.remove(merged_data$GLOB), frequency = 365)
# Perform STL Decomposition
stl_glob <- stlplus(ts_glob, period = 365,
                   s.window = stl_parameters$s.window,
                   t.window = stl_parameters$t.window,
                   l.window = stl_parameters$l.window,
                   robust = stl_parameters$robust)
# Smooth the trend-cycle (using moving average, same order)
smoothed_trendcycle_glob <- ma(stl_glob$data[, "trend"], order = ma_order)
# Extract components (same approach as consumption)
seasonal_glob <- stl_glob$data[, "seasonal"]
deseasoned_glob <- ts_glob - seasonal_glob
trendcycle_glob <- stl_glob$data[, "trend"]
if ("cycle" %in% colnames(stl_glob$data)) {
  trendcycle_glob <- trendcycle_glob + stl_glob$data[, "cycle"]
}
remainder_glob <- stl_glob$data[, "remainder"]
#Plot components
plot(ts_glob, main = "Global Irradiation")
```

## Global Irradiation




```
plot(seasonal_glob, main = "Seasonal Component")
```




```
plot(deseasoned_glob, main = "Deseasoned Irradiation")
```

## Deseasoned Irradiation




```
plot(smoothed_trendcycle_glob, main = "Smoothed Trend-Cycle") # Plot smoothed trend
```

## Smoothed Trend-Cycle




```
plot(remainder_glob, main = "Remainder")
```




```
# --- Global Irradiation ---
ts_glob <- ts(na.remove(merged_data$GLOB), frequency = 365)
# Perform STL Decomposition
stl_glob <- stlplus(ts_glob, period = 365,
  s.window = stl_parameters$s.window,
  t.window = stl_parameters$t.window,
  l.window = stl_parameters$l.window,
  robust = stl_parameters$robust)
# Smooth the trend-cycle (using moving average, same order)
smoothed_trendcycle_glob <- ma(stl_glob$data[, "trend"], order = ma_order)
# Extract components (same approach as consumption)
seasonal_glob <- stl_glob$data[, "seasonal"]
deseasoned_glob <- ts_glob - seasonal_glob
trendcycle_glob <- stl_glob$data[, "trend"]
if ("cycle" %in% colnames(stl_glob$data)) {
  trendcycle_glob <- trendcycle_glob + stl_glob$data[, "cycle"]
}
remainder_glob <- stl_glob$data[, "remainder"]
#Plot components
plot(ts_glob, main = "Global Irradiation")
```



## Global Irradiation




```
plot(seasonal_glob, main = "Seasonal Component")
```




```
plot(deseasoned_glob, main = "Deseasoned Irradiation")
```

## Deseasoned Irradiation




```
plot(smoothed_trendcycle_glob, main = "Smoothed Trend-Cycle") # Plot smoothed trend
```

## Smoothed Trend-Cycle




```
plot(remainder_glob, main = "Remainder")
```




```
# --- Air Temperature ---
# ... (same structure as Global Irradiation) ...
ts_temp <- ts(na.remove(merged_data$LT), frequency = 365)
# Perform STL Decomposition
stl_temp <- stlplus(ts_temp, period = 365,
                    s.window = stl_parameters$s.window,
                    t.window = stl_parameters$t.window,
                    l.window = stl_parameters$l.window,
                    robust = stl_parameters$robust)
# Smooth the trend-cycle (using moving average, same order as others)
smoothed_trendcycle_temp <- ma(stl_temp$data[, "trend"], order = ma_order)
seasonal_temp <- stl_temp$data[, "seasonal"]
deseasoned_temp <- ts_temp - seasonal_temp
trendcycle_temp <- stl_temp$data[, "trend"]
if ("cycle" %in% colnames(stl_temp$data)) {
  trendcycle_temp <- trendcycle_temp + stl_temp$data[, "cycle"]
}
remainder_temp <- stl_temp$data[, "remainder"]
plot(ts_temp, main = "Air Temperature")
```

## Air Temperature




```
plot(seasonal_temp, main = "Seasonal Component")
```



```
plot(deseasoned_temp, main = "Deseasoned Temperature")
```


## Deseasoned Temperature




```
plot(smoothed_trendcycle_temp, main = "Smoothed Trend-Cycle") # Plot smoothed trend
```




## Smoothed Trend-Cycle




```
plot(remainder_temp, main = "Remainder")
```



```
ts_consumption <- ts(na.remove(merged_data$Forretning), frequency = 365)
stl_consumption <- stlplus(ts_consumption, period = 365, s.window = "periodic", t.window = 365, l.window = 365)
plot(stl_consumption)
```



```
# Remove seasonal component
autoplot(ts_consumption - stl_consumption$data[, "seasonal"])
```



```
# Repeat for Temperature and Global Irradiation...
```

## Task G: Granger Causality Test

### 1. Hypotheses

- Null Hypothesis (H0): x does NOT Granger-cause y. Past values of x do not help in predicting y.
- Alternative Hypothesis (H1): x DOES Granger-cause y. Past values of x provide statistically significant information about the future values of y.

### 2. Test Procedure (using AR and VAR models)

#### a. Univariate Autoregression (AR) Model for y:

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \varepsilon_t$$

Model y only using its own past values (lags).

#### b. Vector Autoregression (VAR) Model for x and y:

$$y_t = c + \phi_1 y_{t-1} + \dots + \phi_p y_{t-p} + \beta_1 x_{t-1} + \dots + \beta_p x_{t-p} + \varepsilon_t$$

Model y using both its own past values AND the past values of x.

- #### c. Comparison: Use an F-test to compare the AR and VAR models. If the VAR model (with x) significantly improves prediction of y compared to the AR model (only y's history), reject the null hypothesis. The improvement in predictive ability is assessed by comparing residual sums of squares between restricted (AR) and unrestricted (VAR) models.

```

# Granger Causality test function
granger_test <- function(y, x, group_name, var_name, maxlag = 5) {
  # Align time series and handle potential NA/length issues:
  # Find the maximum start date and minimum end date between the two time series
  start_date <- max(start(y), start(x))
  end_date <- min(end(y), end(x))
  # If the time series do not overlap, skip the test and show a warning
  if (start_date > end_date) {
    warning(paste("No overlapping dates for", group_name, "and", var_name))
    return(NULL)
  }
  # Window the time series to the overlapping date range
  y <- window(y, start = start_date, end = end_date)
  x <- window(x, start = start_date, end = end_date)
  # CRITICAL CHECK: Ensure time series have enough observations after windowing.
  # The Granger test needs at least 2 observations to perform the test.
  if (length(na.remove(y)) <= 1 || length(na.remove(x)) <= 1) {
    warning(paste("Time series too short for Granger test:", group_name, "or", var_name))
    return(NULL)
  }
  # Combine the two time series into a data frame for the Granger test
  combined_data <- data.frame(y = y, x = x)
  # Display a message indicating the variables being tested
  print(paste("Granger Causality Test:", var_name, "->", group_name))
  # Perform the Granger Causality Test using the grangertest() function from the lmtest package
  test_result <- tryCatch({
    # Run the Granger Causality Test with a specified maximum lag (default is 5)
    granger_result <- grangertest(y ~ x, order = maxlag, data = combined_data)
    return(granger_result)
  }, error = function(e) {
    # If an error occurs (e.g., due to data issues), catch and display it
    cat("Error in Granger test:", e, "\n")
    return(NULL)
  })
  # Print the results of the Granger Causality test if successful
  if (!is.null(test_result)) {
    print(test_result)
  } else {
    # If no result, show a message indicating no result was generated
    print(paste("No result for", var_name, "->", group_name))
  }
  # Return the test result (if any)
  return(test_result)
}

```

### Example: Granger Causality Test between LT and Temperature

```

# Extract the consumption data for the specific group (e.g., 'Privat')
consumer_groups <- c("Privat", "Forretning", "Industri")
for (group in consumer_groups) {
  consumption_data <- merged_data[, group]
  ts_consumption <- ts(na.remove(consumption_data), frequency = 365) # Convert to time series with dai

```

```

# Extract temperature (LT) data, assuming it is available in merged_data
ts_temp <- ts(na.remove(merged_data$LT), frequency = 365) # Convert to time series with daily frequency

# Perform the Granger Causality test for Consumption vs Temperature
granger_result <- granger_test(ts_consumption, ts_temp, group, "Temperature (LT)")

# Extract Global Irradiation (GLOB) data
ts_glob <- ts(na.remove(merged_data$GLOB), frequency = 365)

# Perform the Granger Causality test for Consumption vs Global Irradiation
granger_result <- granger_test(ts_consumption, ts_glob, group, "Global Irradiation (GLOB)")

# Perform the Granger Causality test for Temperature vs Global Irradiation
granger_result <- granger_test(ts_temp, ts_glob, group, "Temperature (LT)")
}

## [1] "Granger Causality Test: Temperature (LT) -> Privat"
## [1] "Granger Causality Test: Global Irradiation (GLOB) -> Privat"
## [1] "Granger Causality Test: Temperature (LT) -> Privat"
## [1] "Granger Causality Test: Temperature (LT) -> Forretning"
## [1] "Granger Causality Test: Global Irradiation (GLOB) -> Forretning"
## [1] "Granger Causality Test: Temperature (LT) -> Forretning"
## [1] "Granger Causality Test: Temperature (LT) -> Industri"
## [1] "Granger Causality Test: Global Irradiation (GLOB) -> Industri"
## [1] "Granger Causality Test: Temperature (LT) -> Industri"

```

## Task H: Forecasting with ARIMA Model

### ARIMA Model Forecasting function

```

arima_forecast <- function(time_series, max_order = 5, h = 10) {
  # Ensure the time series is a univariate time series (vector)
  if (is.null(time_series) || length(time_series) < 2) {
    stop("Time series is too short for forecasting")
  }

  # Check for missing data and handle it
  time_series <- na.remove(time_series) # Remove NA values from the series

  # Fit an ARIMA model with automatic order selection
  # auto.arima function automatically selects the best ARIMA model based on AICc criterion
  model <- auto.arima(time_series, max.p = max_order, max.q = max_order, seasonal = FALSE)

  # Display the selected ARIMA model
  print(paste("Fitted ARIMA Model:", as.character(model)))

  # Forecast the next 'h' periods (e.g., 10 days)
  forecast_result <- forecast(model, h = h)

  # Plot the forecasted values along with the historical data
  plot(forecast_result)
  title(main = paste("ARIMA Forecast for next", h, "periods"))

  # Return the forecast results (point forecast, lower and upper prediction intervals)
}

```

```

    return(forecast_result)
}

```

### Example: Forecasting Consumption for a Specific Group (Privat)


```

# Let's forecast the consumption for the 'Privat' group over the next 10 periods
# Extract consumption data for the 'Privat' group from the merged dataset
consumption_data <- merged_data$Privat
# Convert consumption data to a time series object, handling NAs
ts_consumption <- ts(na.remove(consumption_data), frequency = 365) # Daily frequency (365 days per year)
# Apply ARIMA model forecasting for the next 10 periods (e.g., days or time units)
forecast_result <- arima_forecast(ts_consumption, max_order = 5, h = 10)

```

```
## [1] "Fitted ARIMA Model: ARIMA(2,1,1)"
```

### ARIMA Forecast for ARIMA(2,1,1)



### Example: Forecasting Temperature (LT)


```

# Now, let's forecast the temperature (LT) for the next 10 periods
# Extract temperature data from merged dataset
temperature_data <- merged_data$LT
# Convert temperature data to a time series object
ts_temperature <- ts(na.remove(temperature_data), frequency = 365)
# Apply ARIMA model forecasting for the next 10 periods (e.g., days or time units)
forecast_result_temp <- arima_forecast(ts_temperature, max_order = 5, h = 10)

```

```
## [1] "Fitted ARIMA Model: ARIMA(1,1,2)"
```

## ARIMA Forecasts for ARIMA(0,0,2)




### Example: Forecasting Global Irradiation (GLOB)

```
# Similarly, forecast Global Irradiation (GLOB) for the next 10 periods
# Extract global irradiation data from merged dataset
glob_data <- merged_data$GLOB
# Convert global irradiation data to a time series object
ts_glob <- ts(na.remove(glob_data), frequency = 365)
# Apply ARIMA model forecasting for the next 10 periods (e.g., days or time units)
forecast_result_glob <- arima_forecast(ts_glob, max_order = 5, h = 10)
```

```
## [1] "Fitted ARIMA Model: ARIMA(5,0,0) with non-zero mean"
```



## Forecast ARIMA(5,0,0) with 0 periods mean



### Task I: Evaluate Forecast Accuracy for ARIMA Model

```
# Evaluation function to calculate forecast accuracy
evaluate_forecast_accuracy <- function(forecast_result, actual_data) {
  # Ensure that both forecast and actual data are provided and have the same length
  if (length(forecast_result$mean) != length(actual_data)) {
    stop("The forecast and actual data must have the same length")
  }
  # Calculate accuracy measures: MAE, MSE, RMSE, MAPE, etc.
  accuracy_measures <- accuracy(forecast_result, actual_data)
  # Print out the accuracy measures
  print("Accuracy Measures:")
  print(accuracy_measures)
  # Return the accuracy measures for further use
  return(accuracy_measures)
}
```

### Example 1: Evaluating Forecast Accuracy for Consumption (VOLUM\_KWH)

```
# Get the actual consumption data for the last 10 periods
# Assuming that actual future consumption data is available for comparison
actual_consumption <- merged_data$Privat[(length(merged_data$Privat)-9):length(merged_data$Privat)]
# Compare with the forecasted consumption
# Use the forecasted result from Task H
forecast_accuracy_consumption <- evaluate_forecast_accuracy(forecast_result, actual_consumption)
```

```
## [1] "Accuracy Measures:"
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set   -39.88407 19696.33 13877.23  -0.1784362  3.931527 0.9876033
## Test set      -32581.73653 44362.18 38091.36 -12.5854245 14.275539 2.7108558
##           ACF1
## Training set 0.00248148
## Test set      NA
```

### Example 2: Evaluating Forecast Accuracy for Temperature (LT)

```
# Get the actual temperature data for the last 10 periods
actual_temperature <- merged_data$LT[(length(merged_data$LT)-9):length(merged_data$LT)]
# Compare with the forecasted temperature
# Use the forecasted result from Task H
forecast_accuracy_temperature <- evaluate_forecast_accuracy(forecast_result_temp, actual_temperature)
```

```
## [1] "Accuracy Measures:"
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.006203914 2.151897 1.618597 -103.759734 189.2093 0.9835212
## Test set     1.860154277 3.808334 3.582179   8.581756  33.6652 2.1766677
##           ACF1
## Training set 4.783413e-05
## Test set      NA
```

### Example 3: Evaluating Forecast Accuracy for Global Irradiation (GLOB)

```
# Get the actual global irradiation data for the last 10 periods
actual_glob <- merged_data$GLOB[(length(merged_data$GLOB)-9):length(merged_data$GLOB)]
# Compare with the forecasted global irradiation
# Use the forecasted result from Task H
forecast_accuracy_glob <- evaluate_forecast_accuracy(forecast_result_glob, actual_glob)
```

```
## [1] "Accuracy Measures:"
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.01252225 4.822953 3.443017 -56.68061 77.53667 0.9358132
## Test set     -2.77529483 4.109704 3.359272 -75.95699 80.94471 0.9130512
##           ACF1
## Training set -0.02502515
## Test set      NA
```